

# Migrating your Oracle database to Amazon Aurora using DMS and SCT

Arun Thiagarajan – DMS Database Engineer

15<sup>th</sup> June 2018

# Agenda

# What to expect from this session

1. Introduction
2. Migration Playbook
3. Common Oracle to Aurora PostgreSQL Feature Comparison with Examples
4. Schema Conversion Tool

# Introduction - Database Migration Service (DMS) & Schema Conversion Tool (SCT)

# What are DMS and SCT?

**AWS Database Migration Service (DMS)** easily and securely migrates and/or replicates your databases and data warehouses to AWS



**AWS Schema Conversion Tool (SCT)** converts your commercial database and data warehouse schemas to open-source engines or AWS-native services, such as Amazon Aurora and Redshift

# Database Migration Process

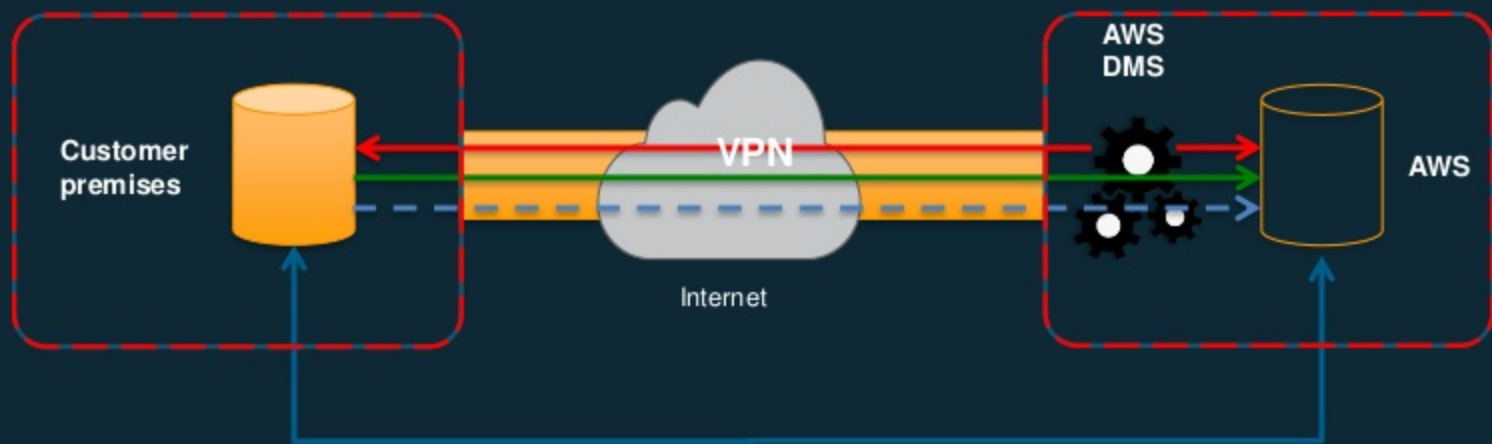
## Step 1: Convert or Copy your Schema



## Step 2: Move your data



# How does DMS work?



Start a replication instance  
Connect to source and target  
databases  
Create Tasks



Application users

- ◆ Let AWS DMS create tables, load data, and keep them in sync
- ◆ Switch applications over to the target at your convenience



# Oracle to Aurora Migration Playbook



# Oracle to Aurora Migration Playbook

- Topic-by-topic overview of Oracle to Aurora PostgreSQL migrations and “hand-on” best practices
- How to migrate from proprietary features and the different database objects
- Migration best practices

Schema



SCT



Data



DMS



Best Practices



Playbook

Oracle Feature	PostgreSQL Feature	Compatibility
<a href="#">Link</a> Index Organized Tables (IOTs)	PostgreSQL “Cluster” Tables	Yes*
<a href="#">Link</a> Common Data Types	Common Data Types	Yes
<a href="#">Link</a> Table Constraints	Table Constraints	Yes
<a href="#">Link</a> Table Partitioning including: RANGE, LIST, HASH, COMPOSITE, Automatic LIST	Table Partitioning including: RANGE, LIST	Yes*
<a href="#">Link</a> Exchange & Split Partitions	N/A	None
<a href="#">Link</a> Temporary Tables	Temporary Tables	Yes*
<a href="#">Link</a> Unused Columns	ALTER TABLE DROP COLUMN	Yes
<a href="#">Link</a> Virtual Columns	Views and/or Function as a Column	Yes*
<a href="#">Link</a> User Defined Types (UDTs)	User Defined Types (UDTs)	Yes
<a href="#">Link</a> Read Only Tables & Table Partitions	Read Only Roles and/or Triggers	Yes*
<a href="#">Link</a> Index Types	Index Types	Yes*
<a href="#">Link</a> B-Trees	Recovery Manager (RMAN)	AWS Aurora Snapshots
<a href="#">Link</a> Comp	Flashback Database	AWS Aurora Snapshots
<a href="#">Link</a> BITM	12c Multi-tenant architecture: PDBs and CDB	Databases
<a href="#">Link</a> Funct	Tablespaces & DataFiles	Tablespaces
<a href="#">Link</a> Globa	Data Pump	pg_dump & pg_restore
<a href="#">Link</a> Index	Resource Manager	Separate AWS Aurora Clusters
<a href="#">Link</a> Ident	Database Users	Database Roles
<a href="#">Link</a> MVCC	Database Roles	Database Roles
<a href="#">Link</a> (Table	SGA & PGA Memory	Memory Buffers
<a href="#">Link</a> Chara	V\$ Views & Data Dictionary	System Catalog Tables, Statistics Collector, AWS Aurora Performance Insights
<a href="#">Link</a> Trans	Log Miner	Logging Options
	Instance & Database Parameters (SPFILE)	AWS Aurora Parameter Groups
	Session Parameters	Session Parameters
	Alert.log (error log)	Error Log via AWS Console
	Automatic and Manual Statistics Collection	Automatic and Manual Statistics Collection
	Viewing Execution Plans	Viewing Execution Plans

# Top Oracle to Aurora PostgreSQL Feature Comparison

# Materialized Views

	Oracle	PostgreSQL
Creation	<pre>CREATE MATERIALIZED VIEW mv1 AS SELECT * FROM employees;</pre>	<pre>CREATE MATERIALIZED VIEW mv1 AS SELECT * FROM employees;</pre>
Manual Refresh	<pre>DBMS_MVIEW.REFRESH('mv1', 'cf');</pre> <p>The --cf parameter configured the refresh method: c is complete and f is fast</p>	<pre>REFRESH MATERIALIZED VIEW mv1;</pre>

# Materialized Views

```
CREATE MATERIALIZED VIEW mv1  
REFRESH FAST ON COMMIT AS  
SELECT * FROM employees;
```

- Supports automatic incremental refresh
- Supports DML on the materialized view

Create a trigger that will initiate a refresh after every DML command on the underlying tables:

```
CREATE OR REPLACE FUNCTION  
refresh_mv1()  
returns trigger language  
plpgsql  
as $$  
begin  
refresh materialized view  
mv1;  
return null;  
end $$;
```

```
create trigger refresh_mv1  
after insert or update or  
delete or truncate  
on employees for each  
statement  
execute procedure
```

# Partitioning

- PostgreSQL started supporting declarative partitioning from version 10.
- Prior to PostgreSQL 10, partitions could be used via inheritance.
- Oracle supports various partitioning mechanisms (hash, range, list, composite) whereas PostgreSQL currently supports list and range partitioning.



# Partitioning

## Oracle example: List Partitioning

```
SQL> CREATE TABLE SYSTEM_LOGS  
  (EVENT_NO NUMBER NOT NULL,  
   EVENT_DATE DATE NOT NULL,  
   EVENT_STR VARCHAR2(500),  
   ERROR_CODE VARCHAR2(10))  
  PARTITION BY LIST (ERROR_CODE)  
  (PARTITION warning VALUES ('err1', 'err2', 'err3') TABLESPACE TB1,  
   PARTITION critical VALUES ('err4', 'err5', 'err6') TABLESPACE TB2);
```

# Partitioning

## PostgreSQL example: List Partitioning

It is a 5 step process using inheritance prior to PostgreSQL 10:

1. Create parent table
2. Create child tables with check constraints
3. Create indexes on child tables
4. Create a function to redirect data inserted into the parent table
5. Create trigger to execute function on DML event



# Partitioning

## 1. Create Parent table:

```
demo=# CREATE TABLE SYSTEM_LOGS  
      (EVENT_NO NUMERIC NOT NULL,  
       EVENT_DATE DATE NOT NULL,  
       EVENT_STR VARCHAR(500),  
       ERROR_CODE VARCHAR(10));
```

## 2. Create Child tables with check constraints:

```
demo=# CREATE TABLE SYSTEM_LOGS_WARNING (  
CHECK (ERROR_CODE IN('err1', 'err2', 'err3')))  
INHERITS (SYSTEM_LOGS);  
demo=# CREATE TABLE SYSTEM_LOGS_CRITICAL (  
CHECK (ERROR_CODE IN('err4', 'err5', 'err6')))  
INHERITS (SYSTEM_LOGS);
```

# Partitioning

3. Create indexes on each of the child tables (“partitions”):

```
demo=# CREATE INDEX IDX_SYSTEM_LOGS_WARNING ON  
SYSTEM_LOGS_WARNING(ERROR_CODE);
```

```
demo=# CREATE INDEX IDX_SYSTEM_LOGS_CRITICAL ON  
SYSTEM_LOGS_CRITICAL(ERROR_CODE);
```

# Partitioning

## 4. Create a function to redirect data inserted into the parent table

```
demo=# CREATE OR REPLACE FUNCTION SYSTEM_LOGS_ERR_CODE_INS()  
      RETURNS TRIGGER AS  
      $$  
      BEGIN  
          IF (NEW.ERROR_CODE IN('err1', 'err2', 'err3')) THEN  
              INSERT INTO SYSTEM_LOGS_WARNING VALUES (NEW.*);  
          ELSIF (NEW.ERROR_CODE IN('err4', 'err5', 'err6')) THEN  
              INSERT INTO SYSTEM_LOGS_CRITICAL VALUES (NEW.*);  
          ELSE  
              RAISE EXCEPTION 'Value out of range, check  
                              SYSTEM_LOGS_ERR_CODE_INS () Function!';  
          END IF;  
          RETURN NULL;  
      END;  
      $$  
      LANGUAGE plpgsql;
```

# Partitioning

## 4. Create trigger to execute function on DML

```
demo=# CREATE TRIGGER SYSTEM_LOGS_ERR_TRIG  
BEFORE INSERT ON SYSTEM_LOGS FOR EACH ROW  
EXECUTE PROCEDURE SYSTEM_LOGS_ERR_CODE_INS();
```

# Triggers

Oracle	PostgreSQL
Triggers can be executed after: <ul style="list-style-type: none"><li>a. DML</li><li>b. DDL</li><li>c. Certain database operations</li></ul>	Triggers can be execute: <ul style="list-style-type: none"><li>a. DML</li><li>b. Event (DDL is also covered in this)</li></ul>
Different types: <ul style="list-style-type: none"><li>a. DML trigger</li><li>b. Instead of trigger</li><li>c. System event trigger</li></ul>	Different types: <ul style="list-style-type: none"><li>a. BEFORE OR AFTER events</li><li>b. INSTEAD OF</li><li>c. For each row or statement</li></ul>

# PostgreSQL triggers

When Fired	Database Event	Row-Level Trigger (FOR EACH ROW)	Statement-Level Trigger (FOR EACH STATEMENT)
BEFORE	INSERT, UPDATE, DELETE	Tables and foreign tables	Tables, views, and foreign tables
	TRUNCATE	—	Tables
AFTER	INSERT, UPDATE, DELETE	Tables and foreign tables	Tables, views, and foreign tables
	TRUNCATE	—	Tables
INSTEAD OF	INSERT, UPDATE, DELETE	Views	—
	TRUNCATE	—	—



# DML Trigger Example - Oracle

```
SQL> CREATE OR REPLACE TRIGGER PROJECTS_SET_NULL
  AFTER DELETE OR UPDATE OF PROJECTNO ON PROJECTS
  FOR EACH ROW
  BEGIN
    IF UPDATING AND :OLD.PROJECTNO != :NEW.PROJECTNO OR DELETING THEN
      UPDATE EMP SET EMP.PROJECTNO = NULL
        WHERE EMP.PROJECTNO = :OLD.PROJECTNO;
    END IF;
  END;
/
```

Trigger created.

```
SQL> DELETE FROM PROJECTS WHERE PROJECTNO=123;
SQL> SELECT PROJECTNO FROM EMP WHERE PROJECTNO=123;
```

```
PROJECTNO
-----
NULL
```



# DML Trigger Example - PostgreSQL

```
psql=> CREATE OR REPLACE FUNCTION PROJECTS_SET_NULL()  
        RETURNS TRIGGER  
        AS $$  
        BEGIN  
            IF TG_OP = 'UPDATE' AND OLD.PROJECTNO != NEW.PROJECTNO OR  
               TG_OP = 'DELETE' THEN  
                UPDATE EMP  
                SET PROJECTNO = NULL  
                WHERE EMP.PROJECTNO = OLD.PROJECTNO;  
            END IF;  
  
            IF TG_OP = 'UPDATE' THEN RETURN NULL;  
            ELSIF TG_OP = 'DELETE' THEN RETURN NULL;  
            END IF;  
        END;  
        $$  
        LANGUAGE PLPGSQL;  
  
CREATE FUNCTION
```

```
psql=> CREATE TRIGGER TRG_PROJECTS_SET_NULL  
        AFTER UPDATE OF PROJECTNO OR DELETE  
        ON PROJECTS  
        FOR EACH ROW  
        EXECUTE PROCEDURE PROJECTS_SET_NULL();  
  
CREATE TRIGGER
```

# Sequences

Parameter/Feature	Compatibility with PostgreSQL	Comments
Create sequence syntax	Full, with minor differences	See Exceptions
INCREMENT BY	Full	
START WITH	Full	
MAXVALUE   NOMAXVALUE	Full	Use "NO MAXVALUE"
MINVALUE   NOMINVALUE	Full	Use "NO MINVALUE"
CYCLE   NOCYCLE	Full	USE "NO CYCLE"
CACHE   NOCACHE	PostgreSQL does not support the NOCACHE parameter but the default behavior is identical. The CACHE parameter is compatible with Oracle.	
Default values using sequences (Oracle 12c)	Supported by PostgreSQL	CREATE TABLE TBL (COL1 NUMERIC DEFAULT NEXTVAL('SEQ 1') ...
Session sequences (session / global), Oracle 12c	Supported by PostgreSQL by using the TEMPORARY sequence parameter to Oracle SESSION sequence	
Oracle 12c identity columns	Supported by PostgreSQL by using the SERIAL data type as sequence	

# Sequence caching in PostgreSQL per session

## In Oracle

```
SQL> create sequence MySeqOracle start with 100 increment by 1 cache 20;  
Sequence created.
```

-Session A

```
SQL> select sys.MySeqOracle.nextval from dual;  
NEXTVAL
```

100

# Open a new session B and generate the next value of sequence

-Session B

```
SQL> select sys.MySeqOracle.nextval from dual;  
NEXTVAL
```

101

# Lets go back to session A and generate the next value

-Session A

```
SQL> select sys.MySeqOracle.nextval from dual;  
NEXTVAL
```

102

## In PostgreSQL

```
psql=> create sequence myseqpostgre start with 100 increment by 1 cache 20;  
CREATE SEQUENCE
```

-Session A

```
psql=> select nextval('public.myseqpostgre');  
nextval
```

100

# Open a new session B and generate the next value of sequence

-Session B

```
psql=> select nextval('public.myseqpostgre');  
nextval
```

120

# Lets go back to session A and generate the next value

-Session A

```
psql=> select nextval('public.myseqpostgre');  
nextval
```

101

# Virtual Columns

Oracle	PostgreSQL
Oracle Virtual Columns appear as normal columns but their values are calculated instead of being stored in the database.	No virtual column support. It can be created through views or function as a column.
Virtual Columns cannot be created based on other Virtual Columns and can only reference columns from the same table	NA
When creating a Virtual Column, you can explicitly specify the datatype or let the database choose the datatype based on the expression.	If the column needs to be part of the object, a view must be created. If not, function as a column can be used as part of the select statement.

# Virtual Columns - Oracle

```
SQL> CREATE TABLE EMPLOYEES (  
    EMPLOYEE_ID NUMBER,  
    FIRST_NAME   VARCHAR2(20),  
    LAST_NAME    VARCHAR2(25),  
    USER_NAME    VARCHAR2(25),  
    EMAIL        AS (LOWER(USER_NAME) || '@aws.com'),  
    HIRE_DATE    DATE,  
    BASE_SALARY  NUMBER,  
    SALES_COUNT  NUMBER,  
    FINEL_SALARY NUMBER GENERATED ALWAYS AS  
    (CASE WHEN SALES_COUNT >= 10 THEN BASE_SALARY + (BASE_SALARY *  
    (SALES_COUNT * 0.05)) END) VIRTUAL);
```



# Virtual Columns - PostgreSQL

1.

```
demo=> CREATE TABLE EMPLOYEES (  
        EMPLOYEE_ID NUMERIC PRIMARY KEY,  
        FIRST_NAME   VARCHAR(20),  
        LAST_NAME    VARCHAR(25),  
        USER_NAME     VARCHAR(25));
```

2.

```
demo=> CREATE OR REPLACE FUNCTION USER_EMAIL(EMPLOYEES)  
        RETURNS text AS $$  
        SELECT (LOWER($1.USER_NAME) || '@aws.com')  
        $$ STABLE LANGUAGE SQL;
```

3.

```
demo=> INSERT INTO EMPLOYEES  
        (EMPLOYEE_ID, FIRST_NAME, LAST_NAME, USER_NAME)  
VALUES (1, 'John', 'Smith', 'jsmith'),  
        (2, 'Steven', 'King', 'sking');
```

# Virtual Columns - PostgreSQL

```
demo=> SELECT EMPLOYEE_ID,  
             FIRST_NAME,  
             LAST_NAME,  
             USER_NAME,  
             USER_EMAIL(EMPLOYEES)  
FROM EMPLOYEES;
```

employee_id	first_name	last_name	user_name	<b>user_email</b>
1	John	Smith	jsmith	<b>jsmith@aws.com</b>
2	Steven	King	sking	<b>sking@aws.com</b>

OR

```
demo=> CREATE VIEW employees_function AS  
SELECT EMPLOYEE_ID,  
       FIRST_NAME,  
       LAST_NAME,  
       USER_NAME,  
       USER_EMAIL(EMPLOYEES)  
FROM EMPLOYEES;
```



# Number in Oracle vs integer or bigint in PostgreSQL

- The limit for numbers in Postgres (up to 131072 digits before the decimal point; up to 16383 digits after the decimal point) is much higher than in Oracle.
- The appropriate data type for NUMBER in Oracle might seem to be NUMERIC in PostgreSQL. However, PostgreSQL Numeric field is less efficient than native integer / bigint fields.
- Given the implementation differences, we have many customers who reduced CPU usage on Aurora from 100% to 20% by changing the generated column data types from Numeric to Bigint on Primary keys.
- If a number field in Oracle is being used as a PK or FK (relating to another table), we should look at using integer or bigint in Aurora PostgreSQL.

# Character Sets

Oracle	PostgreSQL
Supports most national and international character sets including Unicode.	PostgreSQL supports a variety of different character sets, also known as encoding, including support for both single-byte and multi-byte languages.
Supports VARCHAR2 (for non-Unicode) and NVARCHAR2 (for UTF-16).	Does not natively support NVARCHAR2 or UTF-16.
Can be defined at the instance level or pluggable database level.	Called encoding at database level and locale at table/column level.
Changing character set might require an export/import or use database migration assistant for Unicode.	Changing character set requires export/import to a new database.

# Schema Conversion Tool (SCT)

# When to use SCT?

**Modernize** your database tier

ORACLE



## Modernize



**Modernize** and **Migrate** your Data  
Warehouse to Amazon Redshift

VERTICA



TERADATA



ORACLE

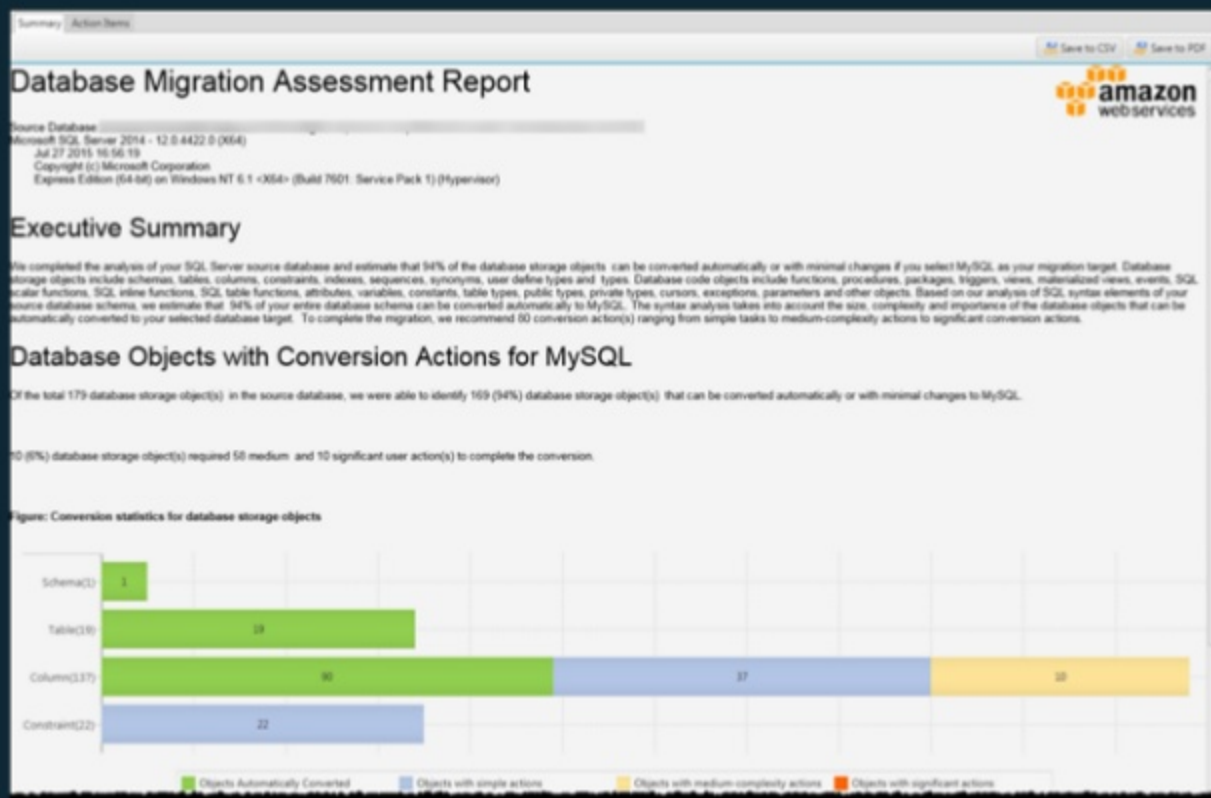


Amazon Redshift



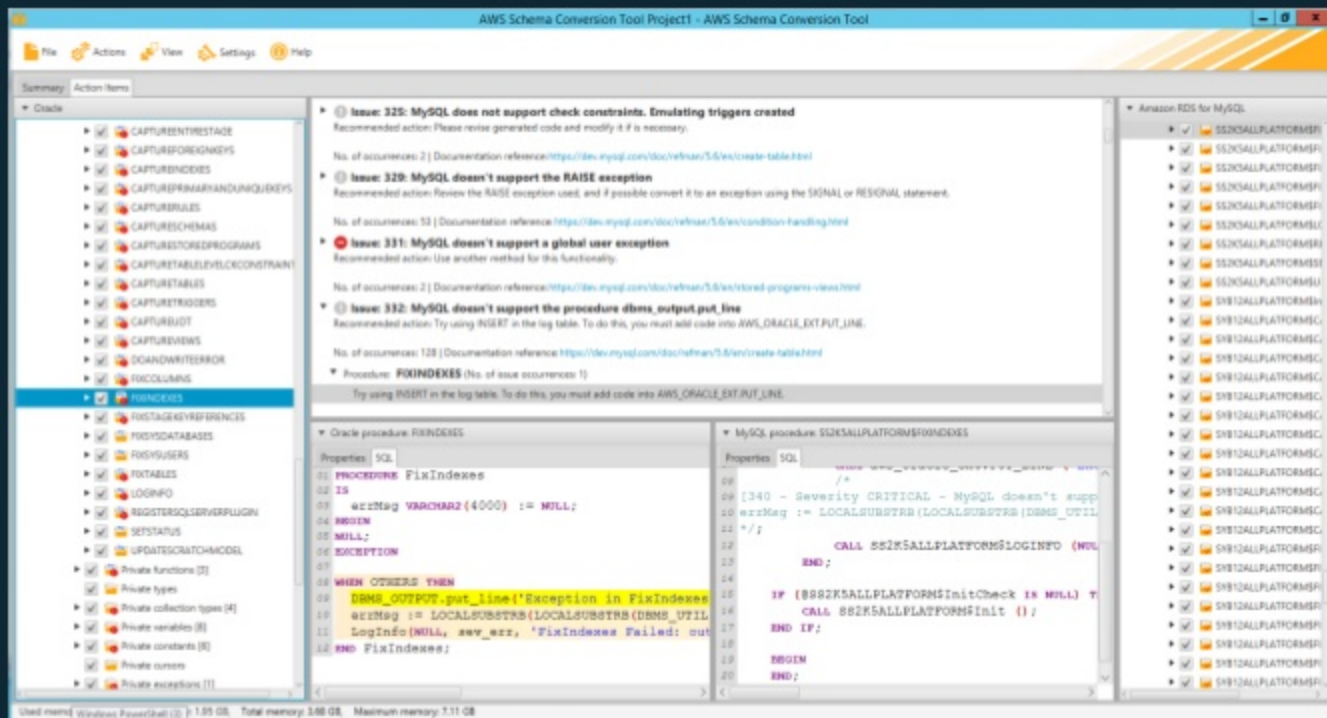


# SCT Migration Assessment Report



- Assessment of migration compatibility of source databases with open-source database engines – RDS MySQL, RDS PostgreSQL and Aurora
- Recommends best target engine
- Provides details level of efforts to complete migration

## SCT helps with converting tables, views, and code



- Sequences
- User-defined types
- Synonyms
- Packages
- Stored procedures
- Functions
- Triggers
- Schemas
- Tables
- Indexes
- Views
- Sort and distribution keys

# Thank you!